

Selection of Open Source Large Language Models

Size	Originator	Largest Model	Parameters	Precision	Size	Description	Parameters	Quantization Level	Size
<= 4 GB	Microsoft	Phi-3 Mini	3.8B	8-bit	~3.8 GB	A surprisingly powerful small model that fits on low-VRAM hardware.	3B	Q4_K_M (4-bit)	~1.8GB
	Alibaba Cloud	Qwen1.5-1.8B-Chat	1.8B	16-bit (FP16)	~3.6 GB	Excellent multilingual model that can run at full native precision.	3B	Q8_0 (8-bit)	~3.1GB
	Google	Gemma 2B	2B	16-bit (FP16)	~4.0 GB	Solid and reliable small model that runs at full precision.	3B	FP16 (unquantized)	~5.8GB
	OpenAI	Cerebras-GPT-2.7B	2.7B	8-bit	~2.7 GB	Good open-source GPT-3 architecture replication for small-scale tasks.	7B	Q8_0 (8-bit)	~7.7GB
6 GB	Meta	Llama 3 8B	8B	4-bit	~4.6 GB	State-of-the-art reasoning and instruction following for small models.	7B	FP16 (unquantized)	~14.0GB
	Microsoft	Phi-3 Small	7B	4-bit	~4.2 GB	Highly capable small model that punches well above its weight.	13B	Q4_K_M (4-bit)	~7.9GB
	Mistral AI	Mistral 7B	7B	4-bit	~4.1 GB	Famous for high efficiency and strong performance for its size.	13B	Q8_0 (8-bit)	~13.8GB
	Alibaba Cloud	Qwen1.5-7B-Chat	7B	4-bit	~4.2 GB	Strong multilingual chat model, great for general purpose conversation.	13B	FP16 (unquantized)	~24.2GB
	Google	Gemma 7B	7B	4-bit	~4.8 GB	Solid, safe, and reliable open model from Google; good all-rounder.	70B	Q4_K_M (4-bit)	~41.4GB
	DeepSeek	DeepSeek LLM 7B	7B	4-bit	~4.1 GB	General purpose base model with good foundational language capabilities.	70B	Q8_0 (8-bit)	~73.3GB
	OpenAI	Cerebras-GPT-6.7B	6.7B	4-bit	~3.8 GB	Open-source replication of GPT-3 architecture, a solid baseline.	70B	FP16 (unquantized)	~140.0GB
	Meta	Llama 2 13B	13B	4-bit	~7.4 GB	A popular, well-supported foundational model for a wide variety of tasks.			
8 GB	Microsoft	Phi-3 Medium	14B	4-bit	~8.1 GB	Powerful medium model with excellent reasoning and language understanding.			
	Mistral AI	Mistral 7B	7B	8-bit	~8.1 GB	Famous for high efficiency and strong performance for its size.			
	Alibaba Cloud	Qwen1.5-14B-Chat	14B	4-bit	~8.2 GB	Larger Qwen model with strong multilingual chat and reasoning.			
	Google	Gemma 7B	7B	8-bit	~8.5 GB	Solid, safe, and reliable open model from Google; good all-rounder.			
	DeepSeek	DeepSeek LLM 7B	7B	8-bit	~7.9 GB	General purpose base model with good foundational language capabilities.			
	OpenAI	Cerebras-GPT-13B	13B	4-bit	~7.4 GB	Larger open-source GPT-3 architecture replication, good general-purpose model.			
	Meta	Llama 2 13B	13B	8-bit	~13.5 GB	A popular, well-supported foundational model for a wide variety of tasks.			
	Microsoft	Phi-3 Medium	14B	4-bit	~8.1 GB	Powerful medium model with excellent reasoning and language understanding.			
12 GB	Mistral AI	Mistral 7B	7B	16-bit (FP16)	~14.1 GB	Famous for high efficiency and strong performance for its size.			
	Alibaba Cloud	Qwen1.5-14B-Chat	14B	4-bit	~8.2 GB	Larger Qwen model with strong multilingual chat and reasoning.			
	Google	Gemma 2 9B	9B	8-bit	~9.1 GB	Google's next-gen small model with leading performance and efficiency.			
	DeepSeek	DeepSeek MoE 16B	16B	4-bit	~9.0 GB	Efficient Mixture-of-Experts model, good at a range of tasks.			
	OpenAI	Cerebras-GPT-13B	13B	8-bit	~13.1 GB	Larger open-source GPT-3 architecture replication, good general-purpose model.			
	Meta	Llama 3 8B	8B	16-bit (FP16)	~16.2 GB	State-of-the-art reasoning and instruction following for small models.			
	Microsoft	Phi-3 Medium	14B	8-bit	~14.1 GB	Powerful medium model with excellent reasoning and language understanding.			
	Mistral AI	Mistral 7B	7B	16-bit (FP16)	~14.1 GB	Famous for high efficiency and strong performance for its size.			
16 GB	Alibaba Cloud	Qwen1.5-14B-Chat	14B	8-bit	~15.2 GB	Larger Qwen model with strong multilingual chat and reasoning.			
	Google	Gemma 2 27B	27B	4-bit	~16.0 GB	Google's powerful next-gen model, excellent for reasoning at scale.			
	DeepSeek	DeepSeek MoE 16B	16B	8-bit	~16.2 GB	Efficient Mixture-of-Experts model, good at a range of tasks.			
	OpenAI	Cerebras-GPT-13B	13B	8-bit	~13.1 GB	Larger open-source GPT-3 architecture replication, good general-purpose model.			
	Meta	Llama 2 13B	13B	16-bit (FP16)	~26.4 GB	A popular, well-supported foundational model for a wide variety of tasks.			
	Microsoft	Phi-3 Medium	14B	16-bit (FP16)	~28.6 GB	Powerful medium model with excellent reasoning and language understanding.			
	Mistral AI	Mistral 7B	7B	16-bit (FP16)	~14.1 GB	Famous for high efficiency and strong performance for its size.			
	Alibaba Cloud	Qwen1.5-32B-Chat	32B	4-bit	~18.5 GB	Very large, powerful multilingual model with strong chat capabilities.			
24 GB	Google	Gemma 2 27B	27B	4-bit	~16.0 GB	Google's powerful next-gen model, excellent for reasoning at scale.			
	DeepSeek	DeepSeek MoE 16B	16B	8-bit	~16.2 GB	Efficient Mixture-of-Experts model, good at a range of tasks.			
	OpenAI	Cerebras-GPT-13B	13B	16-bit (FP16)	~26.2 GB	Larger open-source GPT-3 architecture replication, good general-purpose model.			
	Meta	Llama 2 13B	13B	16-bit (FP16)	~26.4 GB	A popular, well-supported foundational model for a wide variety of tasks.			
	Microsoft	Phi-3 Medium	14B	16-bit (FP16)	~28.6 GB	Powerful medium model with excellent reasoning and language understanding.			
	Mistral AI	Mistral 7B	7B	16-bit (FP16)	~14.1 GB	Famous for high efficiency and strong performance for its size.			
	Alibaba Cloud	Qwen1.5-32B-Chat	32B	4-bit	~18.5 GB	Very large, powerful multilingual model with strong chat capabilities.			
	Google	Gemma 2 27B	27B	8-bit	~27.5 GB	Google's powerful next-gen model, excellent for reasoning at scale.			
32 GB	DeepSeek	DeepSeek MoE 16B	16B	16-bit (FP16)	~32.0 GB	Efficient Mixture-of-Experts model, good at a range of tasks.			
	OpenAI	Cerebras-GPT-13B	13B	16-bit (FP16)	~26.2 GB	Larger open-source GPT-3 architecture replication, good general-purpose model.			
	Meta	Llama 2 13B	13B	16-bit (FP16)	~26.4 GB	A popular, well-supported foundational model for a wide variety of tasks.			
	Microsoft	Phi-3 Medium	14B	16-bit (FP16)	~28.6 GB	Powerful medium model with excellent reasoning and language understanding.			
	Mistral AI	Mixtral 8x7B	47B (MoE)	4-bit	~27.0 GB	Top-tier open-source Mixture-of-Experts model with blazing fast inference.			
	Alibaba Cloud	Qwen1.5-32B-Chat	32B	8-bit	~32.5 GB	Very large, powerful multilingual model with strong chat capabilities.			
	Google	Gemma 2 27B	27B	8-bit	~27.5 GB	Google's powerful next-gen model, excellent for reasoning at scale.			
	DeepSeek	DeepSeek MoE 16B	16B	16-bit (FP16)	~32.0 GB	Efficient Mixture-of-Experts model, good at a range of tasks.			
32 GB	OpenAI	Cerebras-GPT-13B	13B	16-bit (FP16)	~26.2 GB	Larger open-source GPT-3 architecture replication, good general-purpose model.			